

# Learning prompts for transfer learning with test-time adaptation

Trends and Applications of Computer Vision

Giovanni Scialla, Mattia Franzin, Adnan Irshad, Hira Afzal

# Where we left...

Fine-tune the Vision-language model



Improve quality of textual prompts



Prompt Engineering



Prompt Learning



# Where we left...

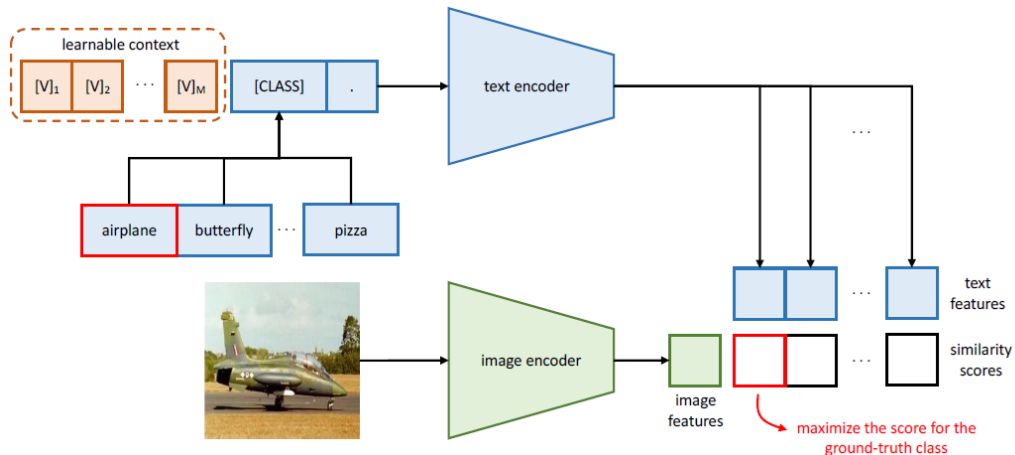
## CoOp

Learn a Context Vector during training

- **Class-specific context**  
one vector for each class
- **Unified context**  
one vector for all classes

```
if csc: # Class-Specific-Context
    ctx_vectors = torch.empty(n_cls, n_ctx, ctx_dim, dtype=dtype)
else: # Generic context
    ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)

self.ctx = nn.Parameter(ctx_vectors) # to be optimized
```



# Where we left...

## CoCoOp

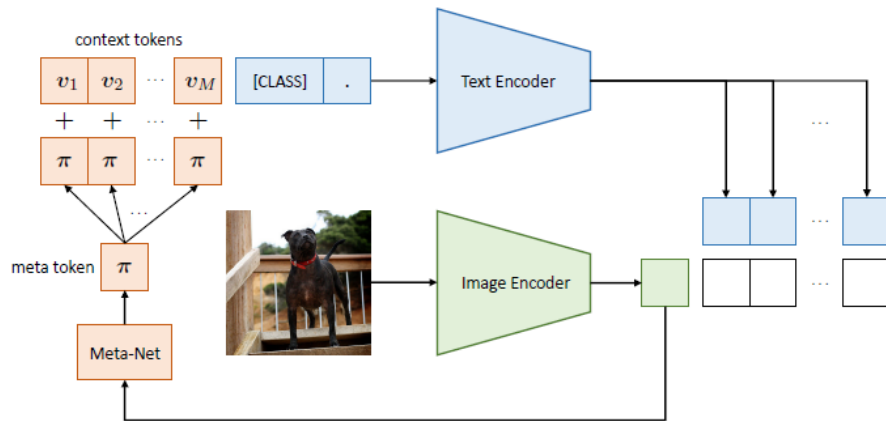
Condition the Context vector using a learnable token from a **MetaNet**

- use image embedding information to update the conditional tokens

```
# random initialization
ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)

self.ctx = nn.Parameter(ctx_vectors)

self.meta_net = nn.Sequential(OrderedDict([
    ('linear1', nn.Linear(vis_dim, vis_dim // 16)),
    ('relu', nn.ReLU(inplace=True)),
    ('linear2', nn.Linear(vis_dim // 16, ctx_dim))
]))
```

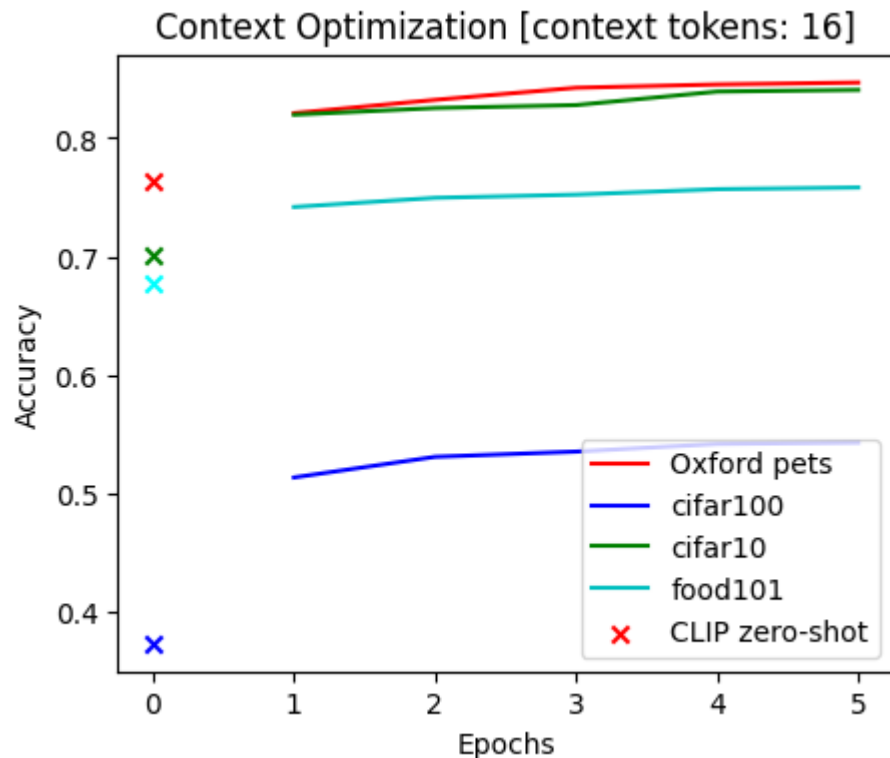


# Datasets

- **CIFAR10** (Alex Krizhevsky, 2009)
- **CIFAR100** (Alex Krizhevsky, 2009)
- **Oxford\_pets** (Parkhi et al., 2012)
- **Food101** (Bossard et al., 2014)
- **MNIST** (Deng, L. 2012)

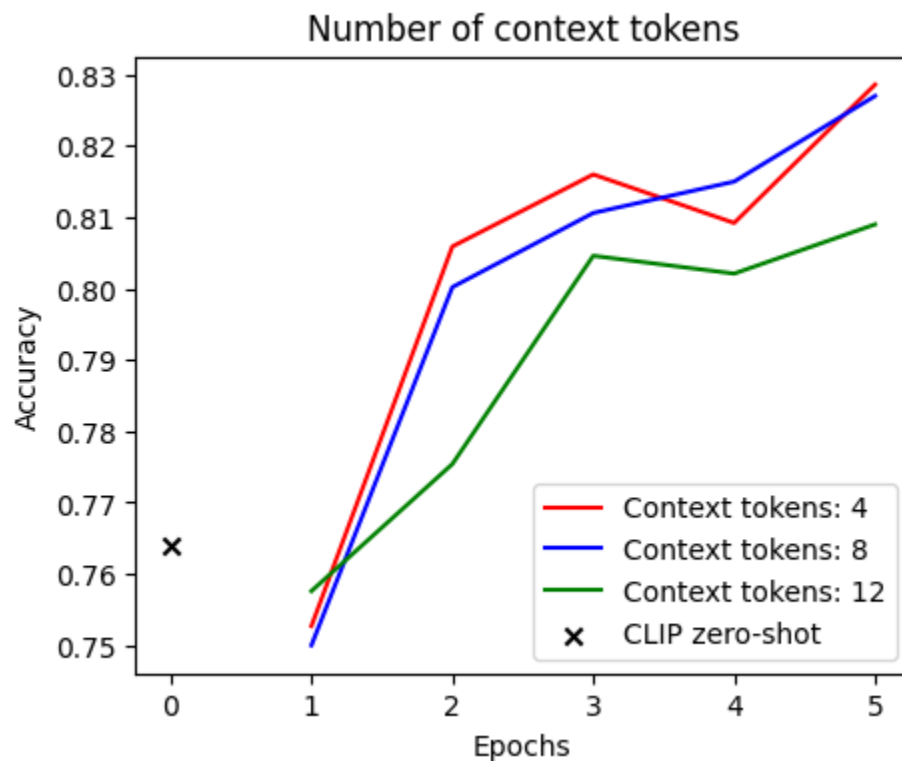
# Experiments and Results

# Context Optimization performance against Zero-shot CLIP



Dataset	Zero-Shot CLIP Accuracy	CoOp Accuracy
Oxford Pets	0.764	0.8468
Food 101	0.677	0.7583
CIFAR10	0.701	0.8406
CIFAR100	0.374	0.5432

# Hyperparameter - Number of Context Tokens








# Comparison between different CoOp modalities

initial context: "X X X X X X X X"


number of context words (tokens): 8




# Dimensionality reduction for visualization - UMAP

**UMAP**: use graph layout algorithms to arrange data to be as structurally similar as possible in a low-dimensional space

CIFAR10 (CLIP Visual Embedding)

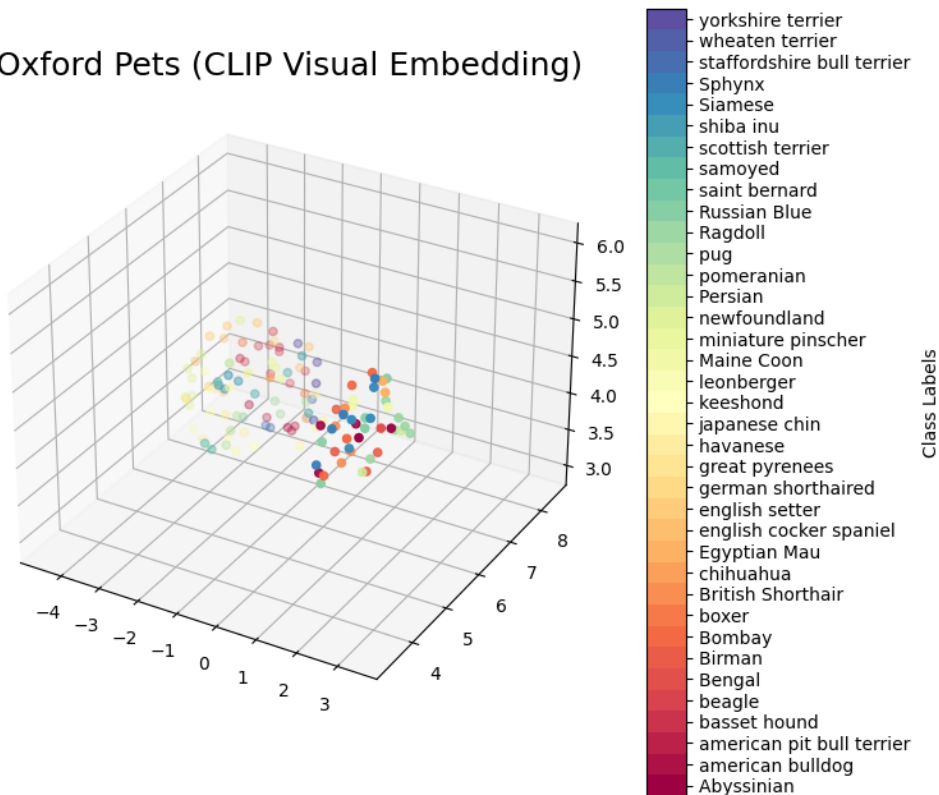


MNIST (CLIP Visual Embedding)



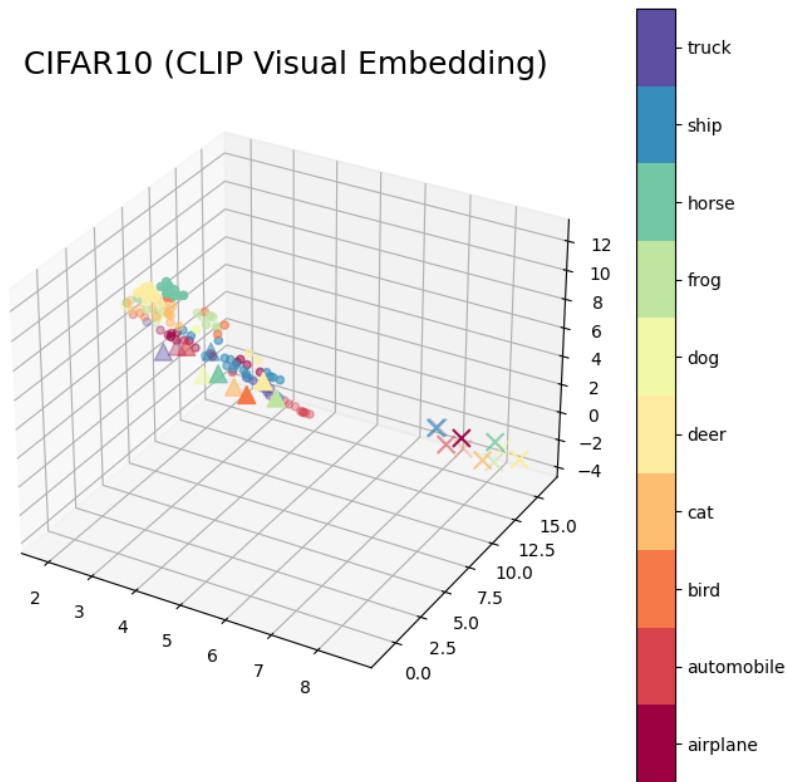
# Dimensionality reduction for visualization - UMAP

Oxford Pets (CLIP Visual Embedding)

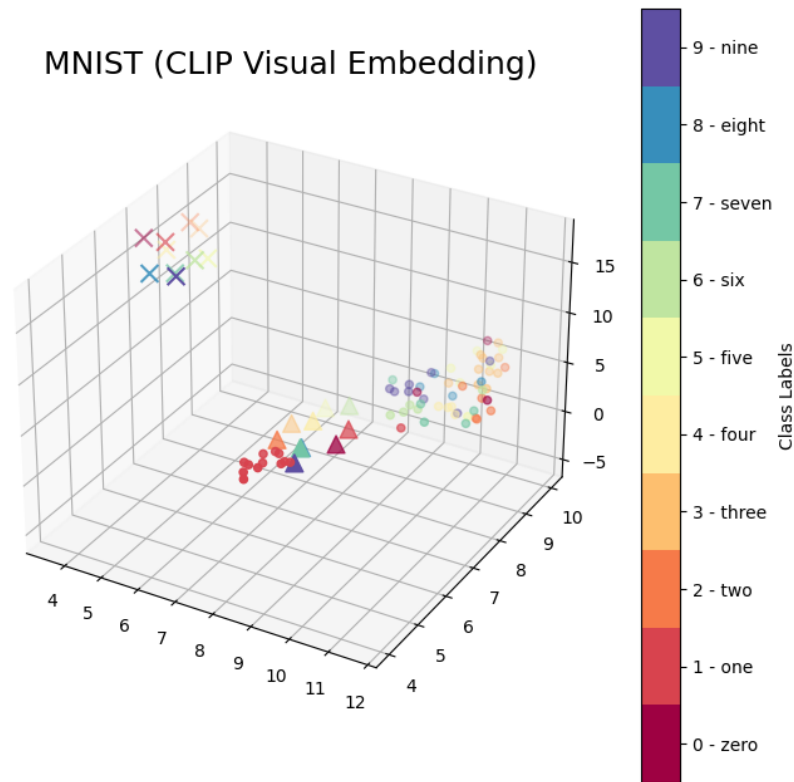


# Learning Prompts for better Embedding alignment

CIFAR10 (CLIP Visual Embedding)

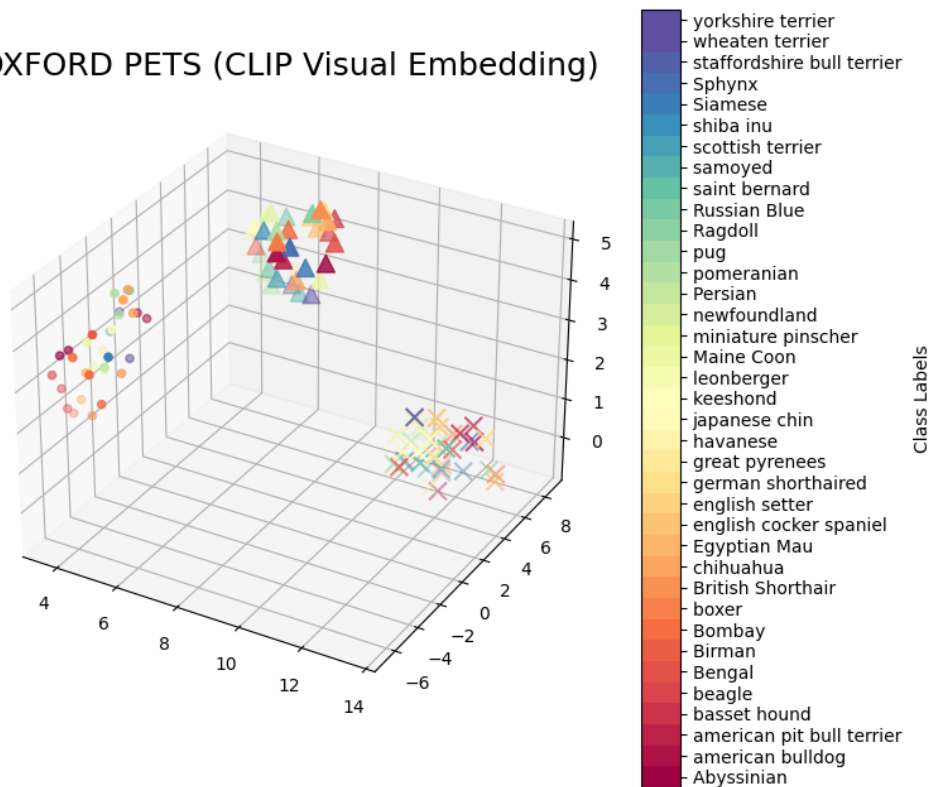


MNIST (CLIP Visual Embedding)



# Learning Prompts for better Embedding alignment

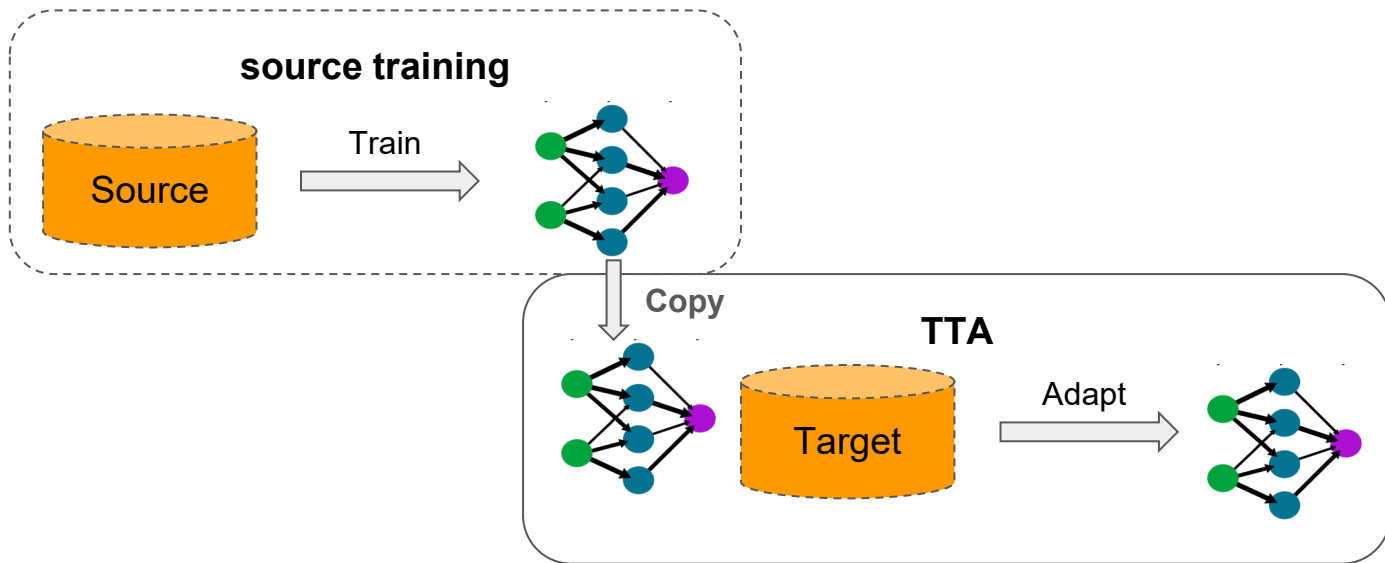
OXFORD PETS (CLIP Visual Embedding)



## **Second Part:** Prompt Tuning with Test-Time Adaption

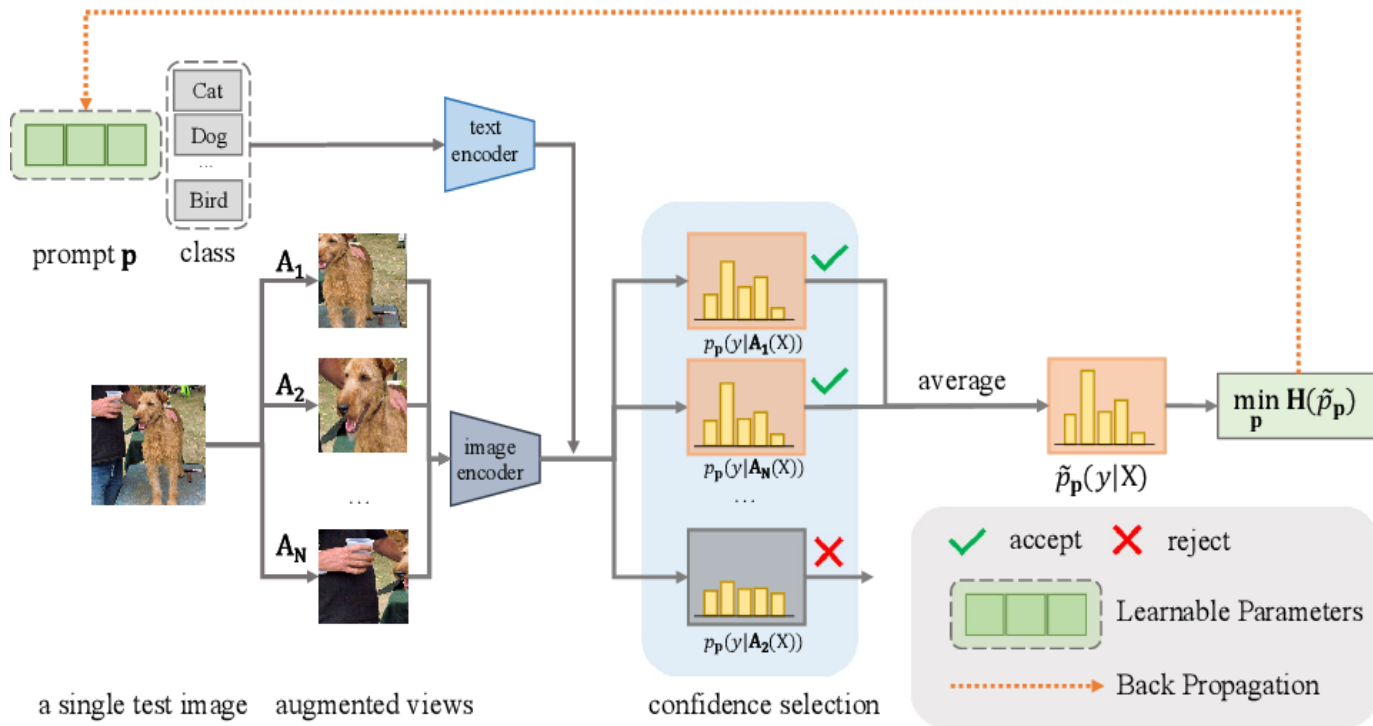
## Second Part: Prompt Tuning with Test-Time Adaption

- Test-Time Prompt Tuning for efficient adaptation
  - An effective lightweight adaptation mechanism at test time for foundation models



# Test-time prompt tuning for zero-shot generalization in Vision Language Models

TPT: learn adaptive prompts  $\mathbf{p}$  on the fly with a single test sample  $\mathbf{X}_{\text{test}}$





# Datasets

- **CIFAR10** (Alex Krizhevsky, 2009)
- **Oxford\_pets** (Parkhi et al., 2012)
- **Food101** (Bossard et al., 2014)
- ImageNet-1K (Jia Deng, 2015)
- ---

BongardHOI (Jiang et al, 2022)

# TPT Baselines

## 1. TPT + CoOP


- a. Image Classification
- b. Context-Dependent Visual Reasoning
  - i. BongardHOI

## 2. DiffTPT + CoOp

- a. Image Classification
  - i. Imagenet-R-1K

## 3. Experiments

- a. Cross-dataset evaluation
  - i. (Oxford\_Pets, Food101, CIFAR10)



# TPT Baselines: TPT + CoOp

## TPT for image classification task

- **Pre-trained CLIP**
  - TextEncoder
  - ImageEncode
- **CoOp**
  - ClipTestTimeTuning
  - PromptLearner
- **AugMix**
  - AugMixAugmenter
- **tpt\_classification**

```
class ClipTestTimeTuning(nn.Module):
    def __init__(self, device, classnames, batch_size,
                 n_ctx=16, ctx_init=None, ctx_posi
    super(ClipTestTimeTuning, self).__init__()
    clip, _, _ = load_arch(device=device, down
    self.image_encoder = clip.visual
    self.text_encoder = TextEncoder(clip)
    self.logit_scale = clip.logit_scale.data

    # prompt tuning
    self.prompt_learner = PromptLearner(
        clip, class_names, batch_size, n_ctx,

    self.criterion = criterion
```


```
def tpt_classification(args):
    model = get_coop(args.arch, args.test_sets, args.gpu, args.n_c

    if args.pretrained_model is not None:
        logger.info("Use pre-trained soft prompt (CoOp) as initial
        pretrained_ctx = torch.load(args.pretrained_model)['state_
        assert pretrained_ctx.size()[0] == args.n_ctx
        with torch.no_grad():
            model.prompt_learner.ctx.copy_(pretrained_ctx)
            model.prompt_learner.ctx_init_state = pretrained_ctx

    model_state = None
```

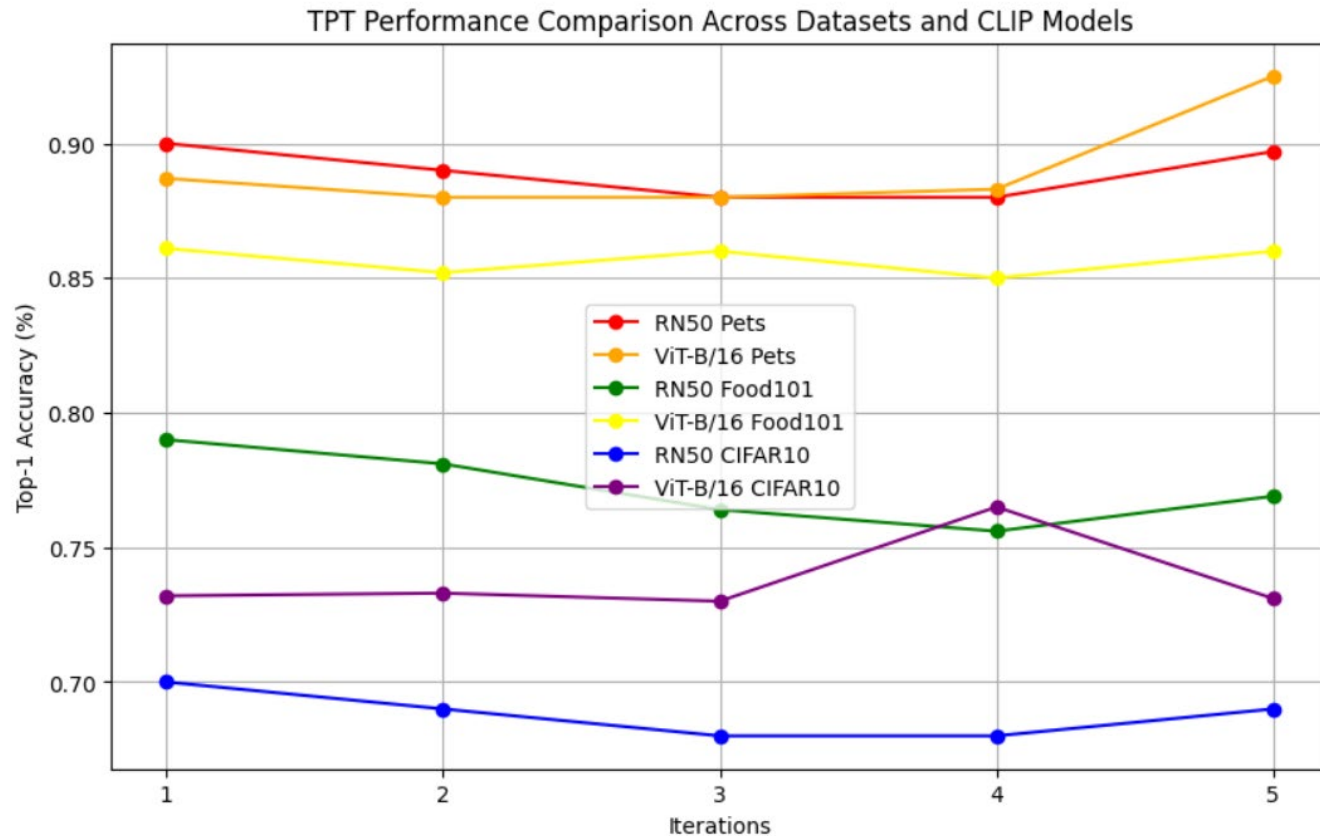
# TPT + CoOP: Performance Analysis

TPT improves by around 5%

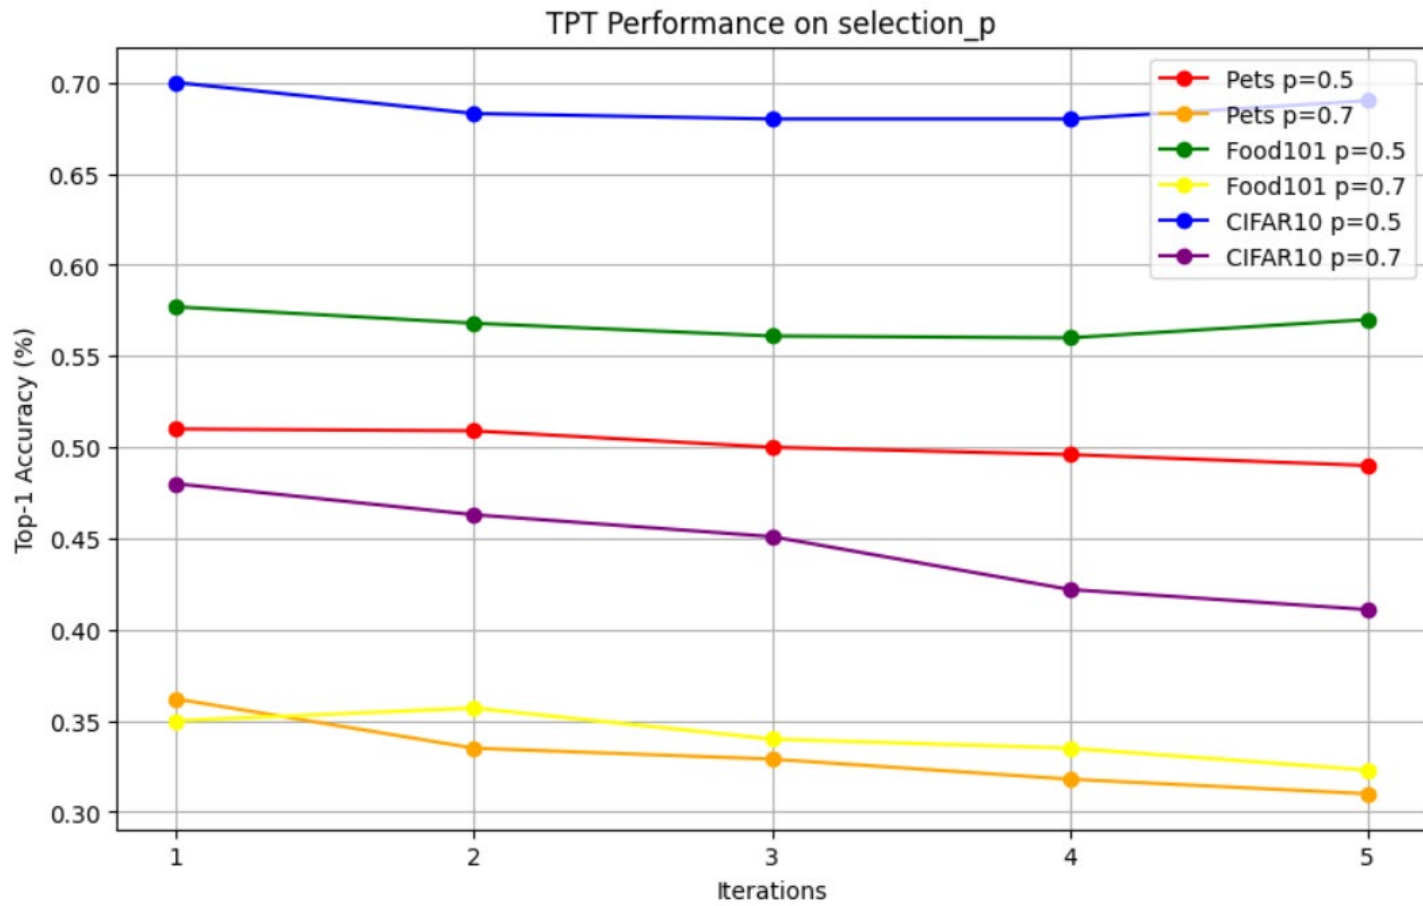


# TPT + CoOP: Performance Comparison with ZSL-CLIP


## ViT-B/16 Performs Well



# Hyperparameters Exp.: Playing with Augmentation



# Hyperparameters Exp.: Playing with Augmentation



# Hyperparameters Exp.: Playing with 'tta\_steps/Iterations'



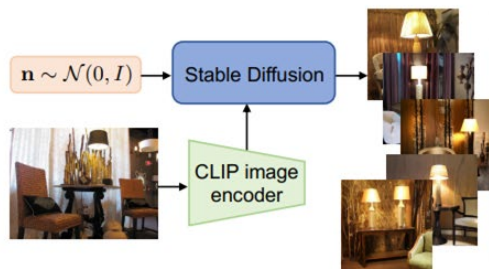


# Thoughts on this TPT Analysis

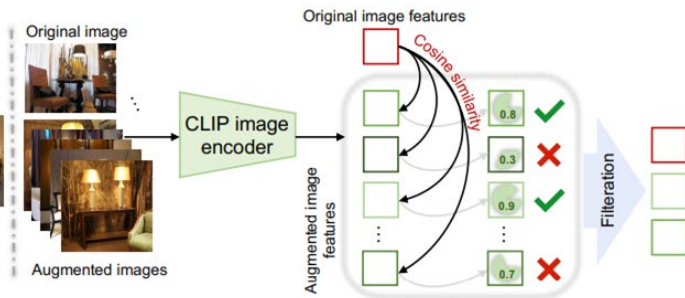
- **Selection Probability ('selection\_p'):** Optimal at 0.1, balances data confidence and performance.
- **Augmented Views ('augmented\_views'):** Best at 64, ensures adaptability without overloading.
- **Test-Time Adaptation Steps ('tta\_steps'):** Ideal at 1, more it, more time it will take and may cause overfitting.
- **Random Seed ('seed'):** Peak performance at 1, crucial for model initialization.
- **Batch Size ('batch\_size'):** Optimal at 64, aligns adaptation capacity with training efficiency.
- Also, we see a major drop in performance during playing with Augmentation especially type/number of augmented views and using selection\_p.

# TPT Baselines: DiffTPT + CoOp

- **Pre-trained CLIP**
  - TextEncoder
  - ImageEncode
- **CoOp**
  - ClipTestTimeTuning
  - PromptLearner
- **AugmentationGenerator**
  - StableDiffusionImageVariationPipeline
- **Diff\_tpt\_classification**




(a) Diverse augmentation with diffusion



(b) Cosine similarity filtration

```
model_name_path = "models/sd-image-variations-diffusers"
pipe = StableDiffusionImageVariationPipeline.from_pretrained(model_name_path,
    torch_dtype=torch.float16)
tform = transforms.Compose([transforms.ToTensor(), transforms.Resize(
    (224, 224), interpolation=transforms.InterpolationMode.BILINEAR,
    antialias=False), transforms.Normalize([0.48145466, 0.4578435, 0.40821056],
    [0.26862954, 0.26130258, 0.27577711])])
dataset_ = DatasetImageNetR(args.data_dir, tform)
dataloader = torch.utils.data.DataLoader(dataset_, batch_size=args.batch_size)
generate_images(pipe, dataloader, args)
```

# DiffTPT Performance over other Models



# Conclusions

We deepened our understanding on prompt learning with test-time adaption by:

- **Studying** several prompt learning methods and test-time adaption techniques drawn from the literature
- **Implementing** from scratch some of those methods
- **Assessing the effectiveness** of the implemented methods by replicating the paper's experiments
- **Visually exploring** in a 3D space the context vector representations of the hand-crafted prompts against the well learned ones